

Student name:	Liu Zihe	CRSID:	zl559	Module number:	3F3
---------------	----------	--------	-------	----------------	-----

Feedback to the student

☐ See also comments in the text

Feedback to the student		Very good	Good	Needs improvmt
C O N T E N T	Completeness, quantity of content: Has the report covered all aspects of the lab? Has the analysis been carried out thoroughly?			
	Correctness, quality of content Is the data correct? Is the analysis of the data correct? Are the conclusions correct?			
	Depth of understanding, quality of discussion Does the report show a good technical understanding? Have all the relevant conclusions been drawn?			
	Comments:			
P R E S E N T A T I O N	Structure, organisation of content Has the report content been structured clearly and in a logical order?			
	Attention to detail, typesetting and typographical errors Is the report free of typographical errors? Are the figures/tables/references presented professionally?			
	Comments:			

Raw report mark	/ 10
Penalty for lateness	

The weighting of comments is not intended to be equal, and the relative importance of criteria may vary between modules. A good report should attract about 7 marks.

2 marks / week or part week.
Please refer to the [online](#) information regarding our extension policy.

Random Variables and Random Number Generation

Liu Zihe^{1*}

¹University of Cambridge

Abstract

This report investigates random number generation and statistical density estimation. We develop theoretical foundations for histogram uncertainty using multinomial distribution theory, implement the inverse CDF transformation method for generating non-uniform random variables, and characterize heavy-tailed α -stable distributions. Empirical validation demonstrates convergence rates for Monte Carlo estimation and the advantage of kernel density estimation over histograms.

1 Introduction

Random number generation underpins modern computational statistics. While pseudorandom number generators readily produce uniform variates, most applications require samples from complex probability distributions. This report develops practical techniques for generating and analyzing non-uniform random variables.

We begin by examining density estimation methods, quantifying histogram uncertainty and how kernel methods achieve better convergence. Probability transformations demonstrate how complex distributions arise from simple ones through functional relationships. The inverse CDF method provides a general framework for sampling arbitrary distributions. Finally, we characterize α -stable distributions, whose heavy tails model phenomena beyond Gaussian assumptions.

2 Uniform and Normal Random Variables

This section compares histogram and kernel density estimation methods, using multinomial distribution theory to quantify histogram uncertainty.

2.1 Kernel Density Estimation

The kernel density estimator is defined as

$$\hat{p}(x) = \frac{1}{N} \sum_{i=1}^N \frac{1}{\sigma} \mathcal{K} \left(\frac{x - x^{(i)}}{\sigma} \right) \quad (1)$$

where N is the number of samples, $x^{(i)}$ are the observed data points, $\mathcal{K}(u)$ is the kernel function, and σ is the bandwidth. The Gaussian kernel is:

$$\mathcal{K}_{\text{Gauss}}(u) = \frac{1}{\sqrt{2\pi}} \exp \left(-\frac{u^2}{2} \right) \quad (2)$$

resulting in the KDE:

$$\hat{p}(x) = \frac{1}{N\sigma} \sum_{i=1}^N \exp \left(-\frac{(x - x^{(i)})^2}{2\sigma^2} \right) \quad (3)$$

2.2 Histogram and KDE of Gaussian Random Numbers

Figure 1 shows that both the histogram (blue bars) and KDE (blue curve) capture the bell-shaped Gaussian profile. The histogram exhibits discontinuities at bin boundaries while the KDE produces a smooth estimate that closely follows the theoretical PDF (red curve).

2.3 Histogram and KDE of Uniform Random Numbers

Figure 2 shows the histogram with approximately equal heights across bins, consistent with uniform density. The KDE ($\sigma = 0.05$) is smoother but extends beyond $[0, 1]$ due to boundary effects.

2.4 Comparison of Methods

Both methods have similar bias-variance structure but differ in asymptotic convergence. From statistical theory [1]: Histogram has an optimal convergence rate of $O(N^{-2/3})$ whereas KDE has an optimal convergence rate of $O(N^{-4/5})$ (faster).

Table 1: Comparison of histogram and KDE methods.

Aspect	Histogram	KDE
Visual Form	Discrete bars	Smooth curve
Interpretation	Bin counts	Smoothed est.
Param. Sens.	High	Lower
Boundary	Respects	May exceed
Convergence	$O(N^{-2/3})$	$O(N^{-4/5})$
Comp. Cost	Trivial	$O(N)$

*CRSID: z1559. Date: 30th November 2025




Figure 1: Normal distribution: histogram (left) and KDE (right) compared with theoretical $\mathcal{N}(0,1)$ PDF. The histogram exhibits discontinuities at bin boundaries while KDE ($\sigma = 0.2$) produces a smooth estimate.




Figure 2: Uniform distribution: histogram (left) and KDE (right) compared with theoretical $\mathcal{U}(0,1)$ PDF. KDE extends beyond $[0,1]$ due to boundary effects.

Histogram advantages: Directly represents sample counts per bin (intuitive), respects finite support of distributions, computationally simple.

Histogram disadvantages: Sensitive to bin width and position, artificial discontinuities between bins.

KDE advantages: Smooth, continuous estimate revealing density structure, superior convergence rate, less sensitive to bandwidth selection.

KDE disadvantages: Less intuitive, boundary effects for finite support distributions, computationally more expensive.

2.5 Multinomial Distribution

The multinomial distribution describes the probability of observing counts $(n_1, n_2, \dots, n_J)$ when N independent samples are distributed into J bins, where each sample falls into bin j with probability p_j :

$$P(n_1, \dots, n_J) = \frac{N!}{n_1! n_2! \dots n_J!} p_1^{n_1} p_2^{n_2} \dots p_J^{n_J} \quad (4)$$

Each sample is an independent trial. For a specific sample to land in bin j , the probability is p_j . If exactly n_j samples land in bin j , these contribute a factor of $p_j^{n_j}$. Since all N samples are independent, the probability of a specific sequence of outcomes is the

product of individual probabilities $p_1^{n_1} \times p_2^{n_2} \times \dots \times p_J^{n_J}$.

The factorial coefficient counts the number of distinct arrangements that yield the same bin counts. Consider N labelled samples: the number of ways to choose which n_1 samples go to bin 1, then which n_2 of the remaining go to bin 2, and so forth, is:

$$\binom{N}{n_1} \binom{N-n_1}{n_2} \dots \binom{n_J}{n_J} = \frac{N!}{n_1! n_2! \dots n_J!} \quad (5)$$

This multinomial coefficient counts the number of ways to partition N distinct objects into J groups of sizes $n_1, \dots, n_J$. Equation (4) thus gives the probability of observing a specific pattern of bin counts.

2.6 Multinomial Analysis of Uniform Distribution

When histogramming N samples from $\mathcal{U}(0,1)$ into J bins, each bin has width $\delta = 1/J$. The probability that a sample falls in bin j is:

$$p_j = \int_{(j-1)/J}^{j/J} p(x) dx = \int_{(j-1)/J}^{j/J} 1 dx = \frac{1}{J} \quad (6)$$

Each marginal count follows a binomial distribution:

$$n_j \sim \text{Binomial} \left(N, p_j = \frac{1}{J} \right) \quad (7)$$

Using the binomial mean formula $E[X] = Np$:

$$E[n_j] = N \cdot p_j = N \cdot \frac{1}{J} = \frac{N}{J} \quad (8)$$

Using the binomial variance formula $\text{Var}(X) = Np(1-p)$:

$$\text{Var}(n_j) = N \cdot p_j \cdot (1 - p_j) = \frac{N(J-1)}{J^2} \quad (9)$$

$$\sigma_{n_j} = \sqrt{\text{Var}(n_j)} = \sqrt{\frac{N(J-1)}{J^2}} = \frac{\sqrt{N(J-1)}}{J} \quad (10)$$

The coefficient of variation (relative uncertainty) is:

$$\begin{aligned} \text{CV} &= \frac{\sigma_{n_j}}{E[n_j]} = \frac{\sqrt{N(J-1)}/J}{N/J} \\ &= \frac{\sqrt{N(J-1)}}{N} = \sqrt{\frac{J-1}{N}} \approx \sqrt{\frac{J}{N}} \end{aligned} \quad (11)$$

As $N \rightarrow \infty$ with J fixed, $\text{CV} \sim 1/\sqrt{N} \rightarrow 0$, so histogram estimates become increasingly reliable. By the Law of Large Numbers, observed bin counts $n_j/N \rightarrow p_j = 1/J$.

Figure 3 also demonstrates good agreement of histogram results with multinomial distribution theory:

1. **Bin Count Distribution:** Expected mean $E[n_j] = N/J$ closely matches observed counts across all N values.
2. **Confidence Bounds:** Most of the bins fall within $\pm 3\sigma$ bounds, consistent with the $\sim 99.7\%$ theoretical expectation from CLT.
3. **Increasing Reliability:** As N increases, relative scatter decreases proportionally to $1/\sqrt{N}$.

2.7 Multinomial Analysis for Gaussian (Normal) Distribution

For normally distributed data $X \sim \mathcal{N}(0,1)$, the bin probabilities p_j are non-uniform. They must be calculated using the CDF:

$$p_j = \Phi(b_{j+1}) - \Phi(b_j) \quad (12)$$

where $\Phi(x)$ is the standard normal CDF and $[b_j, b_{j+1})$ defines bin j . Central bins (near $x = 0$) have high p_j because the PDF is peaked there, while edge bins (large $|x|$) have low p_j due to exponential decay.

Each bin count still follows a binomial distribution:

$$n_j \sim \text{Binomial}(N, p_j) \quad (13)$$

$$E[n_j] = Np_j \quad (14)$$

$$\text{Var}(n_j) = Np_j(1 - p_j) \quad (15)$$

This is maximised when $p_j = 0.5$ and minimised when $p_j \rightarrow 0$ or $p_j \rightarrow 1$.

$$\sigma_{n_j} = \sqrt{Np_j(1 - p_j)} \quad (16)$$

$$\text{CV}_j = \frac{\sigma_{n_j}}{E[n_j]} = \frac{\sqrt{Np_j(1 - p_j)}}{Np_j} = \sqrt{\frac{1 - p_j}{Np_j}} \quad (17)$$

For edge bins where $p_j \rightarrow 0$, we have $(1 - p_j) \rightarrow 1$, so $\text{CV}_j \approx 1/\sqrt{Np_j} \rightarrow \infty$. Thus edge bins have the highest relative uncertainty, making histogram estimates in the tails less reliable. Table 2 illustrates this for bins at 0, 1, 2, and 3 standard deviations from the mean with $N = 1000$ samples. Although absolute variance decreases in the tails, the coefficient of variation increases from 0.07 at the centre to 0.73 at 3 standard deviations.

Table 2: Variance and CV for bins at different distances from the mean ($N = 1000$, bin width = 0.4).

Position	p_j	$E[n_j]$	$\text{Var}(n_j)$	σ_{n_j}	CV
0 sd	0.159	158.5	133.4	11.55	0.07
1 sd	0.097	96.8	87.4	9.35	0.10
2 sd	0.022	22.0	21.5	4.64	0.21
3 sd	0.002	1.9	1.9	1.37	0.73

3 Functions of Random Variables

Random variable transformations reveal how probability distributions change under functional mappings. Using the Jacobian formula, we derive output distributions for linear, quadratic, and sinusoidal transformations.

3.1 Linear Transformation

For $x \sim \mathcal{N}(0,1)$ and $y = ax + b$, the inverse is $x = (y - b)/a$ with Jacobian $|dx/dy| = 1/|a|$. Applying the transformation formula:

$$\begin{aligned} p(y) &= p(x) \left| \frac{dx}{dy} \right| \\ &= \frac{1}{\sqrt{2\pi}} \exp\left(-\frac{(y-b)^2}{2a^2}\right) \cdot \frac{1}{|a|} \\ &= \frac{1}{\sqrt{2\pi a^2}} \exp\left(-\frac{(y-b)^2}{2a^2}\right) \\ &= \mathcal{N}(y|b, a^2) \end{aligned} \quad (18)$$

This demonstrates that a linear transformation of standard normal produces $\mathcal{N}(\mu, \sigma^2)$ with $\mu = b$ and $\sigma^2 = a^2$. Any normal distribution can be generated from $\mathcal{N}(0,1)$ using $Y = \sigma X + \mu$. Figure 5 verifies this result empirically.




Figure 3: Histograms for $N = 100, 1000, 10000$ with theoretical mean (red) and $\pm 3\sigma$ bounds (green). As N increases, relative scatter decreases proportionally to $1/\sqrt{N}$.




Figure 4: Histograms of normal samples for $N = 100, 1000, 10000$ with theoretical mean (red) and $\pm 3\sigma$ bounds (green). Unlike uniform data, the bounds vary across bins due to non-uniform bin probabilities.




Figure 5: Verification of linear transformation $y = 3x + 5$ where $x \sim \mathcal{N}(0,1)$. The histogram matches theoretical $\mathcal{N}(5,9)$.

3.2 Quadratic Transformation

For $x \sim \mathcal{N}(0,1)$ and $y = x^2$ ($y \geq 0$), there are two inverses: $x_1(y) = \sqrt{y}$ and $x_2(y) = -\sqrt{y}$. The Jacobian is $|dx/dy| = 1/(2\sqrt{y})$. Applying the formula:

$$\begin{aligned} p(y) &= 2 \cdot \frac{1}{\sqrt{2\pi}} \exp\left(-\frac{y}{2}\right) \cdot \frac{1}{2\sqrt{y}} \\ &= \frac{1}{\sqrt{2\pi y}} \exp\left(-\frac{y}{2}\right), \quad y \geq 0 \end{aligned} \quad (19)$$

We verify this in Figure 6.




Figure 6: Verification of quadratic transformation $y = x^2$. Sample mean ≈ 1.0 and variance ≈ 2.0

3.3 Sinusoidal Transformation

For $x \sim \mathcal{U}(0, 2\pi)$ and $y = \sin(x)$, the equation $\sin(x) = y$ has two solutions:

$$x_1 = \sin^{-1}(y), \quad x_2 = \pi - \sin^{-1}(y) \quad (20)$$

The Jacobian is:

$$\left| \frac{dx}{dy} \right| = \frac{1}{|\cos(x)|} = \frac{1}{\sqrt{1-y^2}} \quad (21)$$

Applying the multi-valued transformation formula with $p(x) = 1/(2\pi)$:

$$\begin{aligned} p(y) &= 2 \times \frac{1}{2\pi} \times \frac{1}{\sqrt{1-y^2}} \\ &= \frac{1}{\pi\sqrt{1-y^2}}, \quad y \in (-1, 1) \end{aligned} \quad (22)$$

This is the **arcsine distribution**, which has infinite density at $y = \pm 1$. Physically, a sinusoid spends

more time near its peaks and troughs (where the rate of change is slowest) than near zero (where the rate of change is fastest). Figure 7 verifies this result.




Figure 7: Sinusoidal transformation $y = \sin(x)$ where $x \sim \mathcal{U}(0, 2\pi)$. The histogram matches the theoretical arcsine distribution.

3.4 Clipped Sinusoid

Consider the clipped sinusoid $y = \min(\sin(x), 0.7)$, an idealised model of electronic saturation or clipping. The transformation has two regions:

Region 1: No clipping ($\sin(x) < 0.7$):

$$p(y) = \frac{1}{\pi\sqrt{1-y^2}}, \quad -1 < y < 0.7 \quad (23)$$

Region 2: Clipping ($\sin(x) \geq 0.7$): All x values where $\sin(x) \geq 0.7$ map to $y = 0.7$, creating a point mass. The clipped interval $[\sin^{-1}(0.7), \pi - \sin^{-1}(0.7)]$ has length $\pi - 2\sin^{-1}(0.7)$, giving:

$$P(y = 0.7) = \frac{\pi - 2\sin^{-1}(0.7)}{2\pi} \approx 0.253 \quad (24)$$

The complete PDF is:

$$p(y) = \begin{cases} \frac{1}{\pi\sqrt{1-y^2}} & -1 < y < 0.7 \\ \delta(y - 0.7) \cdot 0.253 & y = 0.7 \end{cases} \quad (25)$$

When $\sin(x) = 0.7$, we have $|dy/dx| = |\cos(x)| \rightarrow \sqrt{1-0.7^2} \approx 0.71$. However, the key insight is that a finite range of x values (approximately 25.3% of the input domain) all map to the single output value $y = 0.7$, concentrating probability mass into a delta function. Figure 8 shows the characteristic spike at $y = 0.7$.

4 Inverse CDF Method

The inverse CDF method provides a universal framework for generating samples from arbitrary distributions by transforming uniform random variables.




Figure 8: Clipped sinusoid $y = \min(\sin(x), 0.7)$. The spike at $y = 0.7$ represents approximately 25.3% of samples where saturation occurs.

4.1 Theory

If $U \sim \mathcal{U}(0, 1)$ and $Y = F^{-1}(U)$ where F is any CDF, then Y has distribution with CDF F .

Proof: For any y ,

$$\begin{aligned} \Pr(Y \leq y) &= \Pr(F^{-1}(U) \leq y) \\ &= \Pr(U \leq F(y)) \\ &= F(y) \end{aligned} \quad (26)$$

where the last step uses $\Pr(U \leq u) = u$ for $U \sim \mathcal{U}(0, 1)$.

Algorithm:

1. Generate $u^{(i)} \sim \mathcal{U}(0, 1)$
2. Compute $y^{(i)} = F^{-1}(u^{(i)})$
3. Then $y^{(i)}$ has distribution with CDF F

4.2 Exponential Distribution

For the exponential distribution with PDF $p(y) = \exp(-y)$ for $y \geq 0$:

CDF:

$$F(y) = \int_0^y \exp(-t) dt = 1 - \exp(-y) \quad (27)$$

Inverse CDF: Solving $u = F(y)$ for y :

$$y = -\ln(1 - u) \quad (28)$$

Thus $F^{-1}(u) = -\ln(1 - u)$ for $u \in (0, 1)$. Since $U \sim \mathcal{U}(0, 1)$ implies $(1 - U) \sim \mathcal{U}(0, 1)$, we can generate samples using $Y = -\ln(U)$ instead.

4.3 Implementation

```
def generate_exponential(n, mean=1.0):
    u = np.random.rand(n)
    return -mean * np.log(u)
```

Figure 9 shows excellent agreement between the histogram and theoretical exponential density. The KDE exhibits boundary effects near $y = 0$ where Gaussian kernels extend beyond the valid support $[0, \infty)$.




Figure 9: Exponential distribution via inverse CDF method. Histogram (top) matches theoretical PDF. KDE (bottom) shows boundary effects: peak underestimated (≈ 0.65 vs 1.0) due to Gaussian kernels extending into $y < 0$.

4.4 Monte Carlo Estimation

Given N i.i.d. samples $y^{(1)}, \dots, y^{(N)}$ from a distribution with mean μ and variance σ^2 , we define the sample mean estimator:

$$\hat{\mu} = \frac{1}{N} \sum_{i=1}^N y^{(i)} \quad (29)$$

For the exponential distribution with $\lambda = 1$, the theoretical mean is $E[Y] = 1$ and variance is $\text{Var}(Y) = 1$. Experimental verification with $N = 10000$ samples yields $\hat{\mu} = 1.021$ and $\hat{\sigma}^2 = 1.025$, confirming the inverse CDF method correctly generates exponential variates.

4.5 Bias of the Mean Estimator

The sample mean $\hat{\mu}$ is an unbiased estimator of the true mean μ . To prove this, we compute the expecta-

tion over all possible samples:

$$\begin{aligned}
 E[\hat{\mu}] &= E\left[\frac{1}{N} \sum_{i=1}^N y^{(i)}\right] \\
 &= \frac{1}{N} \sum_{i=1}^N E[y^{(i)}] \\
 &= \frac{1}{N} \cdot N \cdot \mu = \mu
 \end{aligned} \tag{30}$$

where we used linearity of expectation and that each $y^{(i)}$ has the same mean μ .

4.6 Variance of the Mean Estimator

The variance of $\hat{\mu}$ determines how quickly the estimate converges to the true mean:

$$\begin{aligned}
 \text{Var}(\hat{\mu}) &= \text{Var}\left(\frac{1}{N} \sum_{i=1}^N y^{(i)}\right) \\
 &= \frac{1}{N^2} \text{Var}\left(\sum_{i=1}^N y^{(i)}\right) \\
 &= \frac{1}{N^2} \cdot N \sigma^2 = \frac{\sigma^2}{N}
 \end{aligned} \tag{31}$$

where we used independence of samples (variance of sum equals sum of variances).

The standard error is $\text{SE}(\hat{\mu}) = \sigma/\sqrt{N}$, implying that doubling the number of samples only reduces the error by a factor of $\sqrt{2}$. This $O(1/\sqrt{N})$ convergence rate is fundamental to Monte Carlo methods.

4.7 Monte Carlo Convergence

Figure 10 demonstrates the $1/N$ variance decay. While individual runs show fluctuations, the averaged squared error closely follows the theoretical σ^2/N bound. This confirms that the Monte Carlo estimator is consistent: as $N \rightarrow \infty$, $\hat{\mu} \rightarrow \mu$ by the Law of Large Numbers.

5 Alpha-Stable Distribution

The α -stable family generalizes the Gaussian distribution with power-law tails, providing a model for heavy-tailed phenomena in communications, finance, and signal processing.

5.1 Algorithm

To generate samples with $\alpha \in (0, 2)$, $\alpha \neq 1$ and $\beta \in [-1, +1]$:




Figure 10: Monte Carlo convergence for exponential distribution. Top: Single run showing squared error $(\hat{\mu} - \mu)^2$ versus N with theoretical variance σ^2/N (dashed). Bottom: Averaged squared error over 100 runs confirms $1/N$ decay rate.

1. Calculate auxiliary parameters:

$$b = \frac{1}{\alpha} \tan^{-1} \left(\beta \tan \frac{\pi\alpha}{2} \right) \tag{32}$$

$$s = \left(1 + \beta^2 \tan^2 \frac{\pi\alpha}{2} \right)^{1/(2\alpha)} \tag{33}$$

2. Generate $U \sim \mathcal{U}(-\pi/2, +\pi/2)$ and $V \sim \mathcal{E}(1)$
3. Calculate:

$$X = s \frac{\sin(\alpha(U+b))}{(\cos U)^{1/\alpha}} \left(\frac{\cos(U - \alpha(U+b))}{V} \right)^{\frac{1-\alpha}{\alpha}} \tag{34}$$

5.2 Implementation

```
def generate_stable(alpha, beta, n):
    b = (1/alpha) * np.arctan(
        beta * np.tan(np.pi*alpha/2))
    s = (1 + beta**2 * np.tan(
        np.pi*alpha/2)**2)**(1/(2*alpha))
    U = np.random.uniform(
        -np.pi/2, np.pi/2, n)
    V = np.random.exponential(1.0, n)
    num = np.sin(alpha*(U+b)) / \
        np.cos(U)**(1/alpha)
    den = (np.cos(U-alpha*(U+b))/V) \
        **((1-alpha)/alpha)
```

```
return s * num * den
```

5.3 Parameter Interpretation

Stability parameter $\alpha \in (0, 2)$: Controls tail heaviness. At $\alpha = 0.5$, tails are extremely heavy with massive outliers. At $\alpha = 1.5$, tails are lighter but still heavier than Gaussian. As $\alpha \rightarrow 2$, the distribution approaches Gaussian.

Skewness parameter $\beta \in [-1, +1]$: Controls asymmetry. At $\beta = 0$, the distribution is symmetric. Positive β produces right skewness; negative β produces left skewness. Figure 11 illustrates these effects for different α and β combinations.

5.4 Tail Probability Analysis

A key distinction between α -stable and Gaussian distributions lies in their tail behaviour. For α -stable distributions with $\alpha < 2$, the tail probability follows a power law:

$$P(|X| > t) \sim c \cdot t^{-\alpha} \quad \text{as } t \rightarrow \infty \quad (35)$$

In contrast, the Gaussian has exponentially decaying tails: $P(|Z| > t) \sim e^{-t^2/2} / (t\sqrt{\pi})$.

Table 3: Tail probabilities $P(|X| > t)$ for different distributions.

Distribution	$t = 0$	$t = 3$	$t = 6$
Gaussian ($\alpha = 2$)	1.0	0.0027	2×10^{-9}
α -stable ($\alpha = 1.5$)	1.0	0.11	0.04
α -stable ($\alpha = 0.5$)	1.0	0.30	0.21

Table 3 demonstrates the dramatic difference: at $t = 6$, the Gaussian tail probability is essentially zero ($\sim 10^{-9}$), while the α -stable with $\alpha = 0.5$ still has $P(|X| > 6) \approx 0.21$. This explains why heavy-tailed distributions produce extreme outliers.

Figure 12 confirms that $P(|X| > t) \propto t^\gamma$ with $\gamma \approx -\alpha$. Linear regression on the log-log plot yields $\gamma = -0.50$ for $\alpha = 0.5$ and $\gamma = -1.50$ for $\alpha = 1.5$, validating the general formula $\gamma = -\alpha$.

5.5 Convergence to Gaussian

As the stability parameter α approaches 2, the α -stable distribution converges to a Gaussian. This transition involves fundamental changes in the distribution's properties:

Tail behaviour: For $\alpha < 2$, tails follow a power law with infinite variance. At $\alpha = 2$, tails decay exponentially with finite variance.

Characteristic function: The α -stable characteristic function is $\phi(t) = \exp(-|t|^\alpha)$. As $\alpha \rightarrow 2$, this approaches $\exp(-t^2)$, the Gaussian characteristic function.

Figure 13 illustrates this convergence visually. At $\alpha = 1.99$, the histogram is nearly indistinguishable from the Gaussian reference, while at $\alpha = 0.5$ the heavy tails produce a sharply peaked distribution with frequent extreme values.

6 Conclusion

This work develops a comprehensive analysis on random variables. Beginning with histogram analysis, multinomial distribution theory reveals that estimation uncertainty scales as $\sqrt{J/N}$ for uniform distributions but diverges in low-probability regions, motivating kernel density estimation with its superior convergence rates. The Jacobian transformation formula then shows how functional mappings reshape probability distributions. The inverse CDF method provides universal sampling from arbitrary distributions by inverting the transformation $U = F(Y)$ for uniform U .

The practical limitations of these techniques emerge through variance analysis and tail behavior. Monte Carlo estimation exhibits fundamental $O(1/\sqrt{N})$ convergence, meaning quadrupled sample sizes merely halve estimation error regardless of algorithmic sophistication. While Gaussian tails decay exponentially as $\exp(-t^2/2)$, α -stable tails follow power laws $t^{-\alpha}$, producing extreme values orders of magnitude more frequently. The stability parameter α directly governs this tail heaviness, with the Gaussian recovered as the limiting case $\alpha \rightarrow 2$. These tools collectively enable practitioners to generate, analyze, and understand random phenomena beyond standard distributional assumptions.

References

- [1] University of Washington. *STAT 425: Course Lecture Notes on Density Estimation*. Lecture notes covering histogram and kernel density estimation methods. 2018. URL: https://faculty.washington.edu/yenchic/18W_425/Lec6_hist_KDE.pdf.




Figure 11: Alpha-stable distributions for $\alpha = 0.5$ (top, heavy tails) and $\alpha = 1.5$ (bottom, lighter tails) with varying β (skewness). Heavy-tailed distributions ($\alpha = 0.5$) produce extreme outliers requiring plot clipping.




Figure 12: Power law tail fitting for α -stable distributions. Log-log plots of $P(|X| > t)$ versus t show linear decay with slope $\gamma \approx -\alpha$, confirming the theoretical prediction. As α approaches 2, the log-log plot for the α -stable distribution is also closer to that for the gaussian distribution.




Figure 13: Convergence of α -stable to Gaussian as $\alpha \rightarrow 2$. For $\alpha = 0.5$, the distribution has extremely heavy tails. As α increases through 1.5, 1.7, 1.9, 1.99, the distribution progressively approaches the Gaussian (red curve).